

Key Management and Identity



CS 458: Information Security
Kevin Jin

Administrivia

- Lab 1 due on 9/27
- Homework 3 solution and Homework 4 released

Reading Material

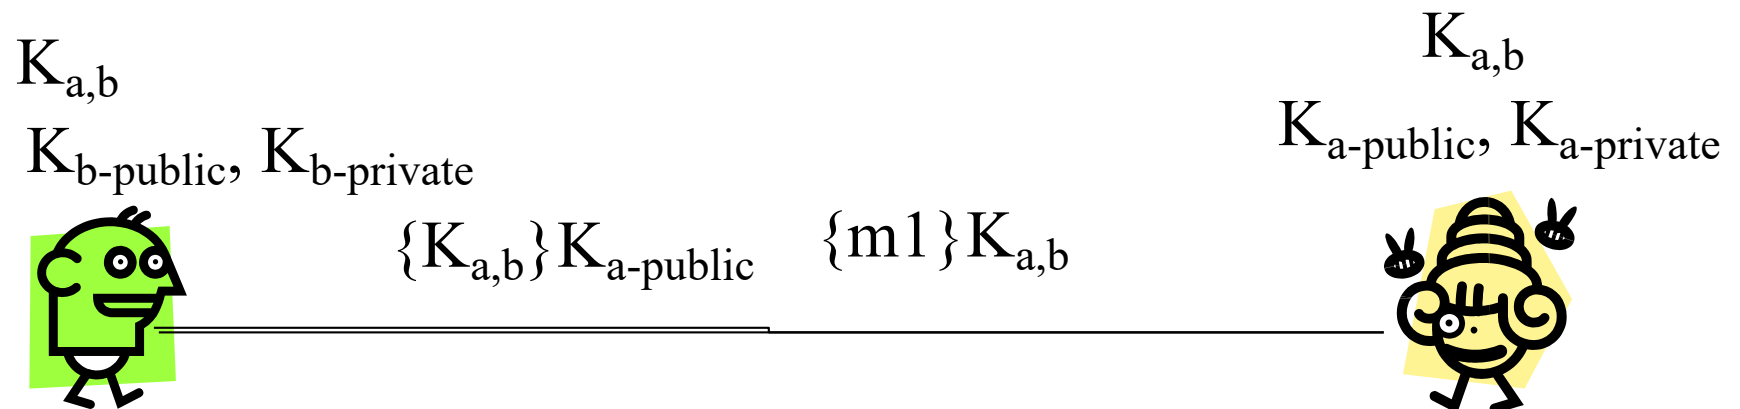
- Chapters 2.4, 2.5 and 23.1 - 3 from text
- Handbook of Applied Cryptography
 - <http://www.cacr.math.uwaterloo.ca/hac/>
 - Chapter 5 for information on Pseudorandom sequences
 - Chapter 12 for Needham-Schroeder protocol
- The Gnu Privacy Handbook
 - <http://www.gnupg.org/gph/en/manual.html>

Overview

- Key Exchange & Random Numbers
- Key/Identity Management
 - Kerberos
 - Digital Signatures & PKI
 - Hierarchical – X.509
 - Web of Trust

Session and Interchange Keys

- Long lived Interchange Keys only exist to boot strap
- Short lived session keys used for bulk encryption



Picking Good Session Keys

- Goal: generate keys that are difficult to guess
- Problem statement: given a set of K potential keys, choose one randomly
 - Equivalent to selecting a random number between 0 and $K-1$ inclusive
 - Uniform distribution (use entire space)
 - Independence (should not be able to predict next selection)
- Why is this hard?
 - Actually, numbers are usually *pseudo-random*, that is, generated by an algorithm

What is “Random”?



- *Sequence of cryptographically random numbers:*
a sequence of numbers n_1, n_2, \dots such that for any integer $k > 0$, an observer cannot predict n_k even if all of n_1, \dots, n_{k-1} are known
 - Best: physical source of randomness
 - Random pulses
 - Electromagnetic phenomena
 - Characteristics of computing environment such as disk latency
 - Ambient background noise

What is “Pseudorandom”?



- *Sequence of cryptographically pseudorandom numbers:*
sequence of numbers intended to simulate a sequence of cryptographically random numbers but **generated by an algorithm**
 - Very difficult to do this well
 - Linear congruential generators
 $[n_k = (an_{k-1} + b) \bmod n]$ broken
 - Polynomial congruential generators
 $[n_k = (a_j n_{k-1}^j + \dots + a_1 n_{k-1} + a_0) \bmod n]$ broken too
 - Here, “broken” means next number in sequence can be determined

Best Pseudorandom Numbers

- *Strong mixing function*: function of 2 or more inputs with each bit of output depending on some nonlinear function of all input bits
 - Examples: DES, MD5, SHA-1, avalanche effect
 - Use on UNIX-based systems:

```
(date; ps aux) | md5
```

where “ps aux” lists all information about all processes on system

Separate Channel

- Ideally you have separate secure channel for exchanging Interchange keys
 - Direct secret sharing grows at N^2

Telephone, separate data network, ESP, sneaker net




Regular data network

So how can we solve that?

Shared Channel: Trusted Third Party


- Generally separate channel is not practical
 - No trustworthy separate channel
 - Want to scale linearly with additional users



Classical Key Exchange

- Bootstrap problem: how do Alice, Bob begin?
 - Alice can't send it to Bob in the clear!
- Assume trusted third party, Cathy
 - Alice and Cathy share secret key k_A
 - Bob and Cathy share secret key k_B
- Use this to exchange shared key k_s

Simple Protocol




What is the problem?

Problems

- How does Bob know he is talking to Alice?
 - Replay attack
 - Eve records message from Alice to Bob, later replays it
 - Bob may think he's talking to Alice, but he isn't
 - Session key reuse
 - Eve replays message from Alice to Bob, so Bob re-uses session key
- Protocols must provide authentication and defence against replay

Needham-Schroeder



Argument: Alice talking to Bob

- Second message
 - Encrypted using key only Alice and Cathy know
 - So Cathy encrypted it
 - Response to first message
 - As r_1 in it matches r_1 in first message

Alice $\xleftarrow{\{\text{Alice}||\text{Bob}||r_1||k_s \ \{\text{Alice}||k_s\} \ k_B\} \ k_A}$ Cathy


Argument: Alice talking to Bob

- Third message
 - Alice knows only Bob can read it
 - As only Bob can derive session key from message
 - Encrypted using key only Bob and Cathy know
 - So Cathy encrypted it
 - Names Alice, session key
 - Cathy provided session key, says Alice is other party

Alice $\xrightarrow{\{Alice || k_s\} k_B}$ Bob

Argument: Bob talking to Alice

- Fourth and fifth messages
 - Uses session key to determine if it is replay from Eve
 - If not, Alice will respond correctly in fifth message
 - If so, Eve can't decrypt r_2 and so can't respond, or responds incorrectly



Kerberos



Kerberos

- In Greek mythology, Kerberos is a three-headed dog that guards entrance to Hades
 - “Wouldn’t it make more sense to guard the exit?”
- In security, Kerberos is an authentication protocol based on symmetric key crypto
 - Originated at MIT
 - Based on work by Needham and Schroeder
 - Relies on a **Trusted Third Party (TTP)**

Motivation for Kerberos

- Authentication using public keys
 - N users \rightarrow N key pairs
- Authentication using symmetric keys
 - N users requires (on the order of) N^2 keys
- Symmetric key case **does not scale**
- Kerberos based on symmetric keys but only requires N keys for N users
 - Security depends on TTP
 - + No PKI is needed

Kerberos KDC

- Kerberos **Key Distribution Center** or **KDC**
 - KDC acts as the TTP
 - TTP is trusted, so it must not be compromised
- KDC shares symmetric key K_A with Alice, key K_B with Bob, key K_C with Carol, etc.
- A master key K_{KDC} known ***only*** to KDC
- KDC enables authentication, session keys
 - Session key for confidentiality and integrity
- In practice, crypto algorithm is DES


Kerberos Tickets

- KDC issue **tickets** containing info needed to access network resources
- KDC also issues **Ticket-Granting Tickets** or **TGTs** that are used to obtain tickets
- Each TGT contains
 - Session key
 - User's ID
 - Expiration time
- Every TGT is encrypted with K_{KDC}
 - So, TGT can only be read by the KDC

Kerberized Login


- Alice enters her password
- Then Alice's computer does following:
 - Derives K_A from Alice's password
 - Uses K_A to get TGT for Alice from KDC
- Alice then uses her TGT (credentials) to securely access network resources
- **Plus:** Security is transparent to Alice
- **Minus:** KDC *must* be secure - it's trusted!

Kerberized Login



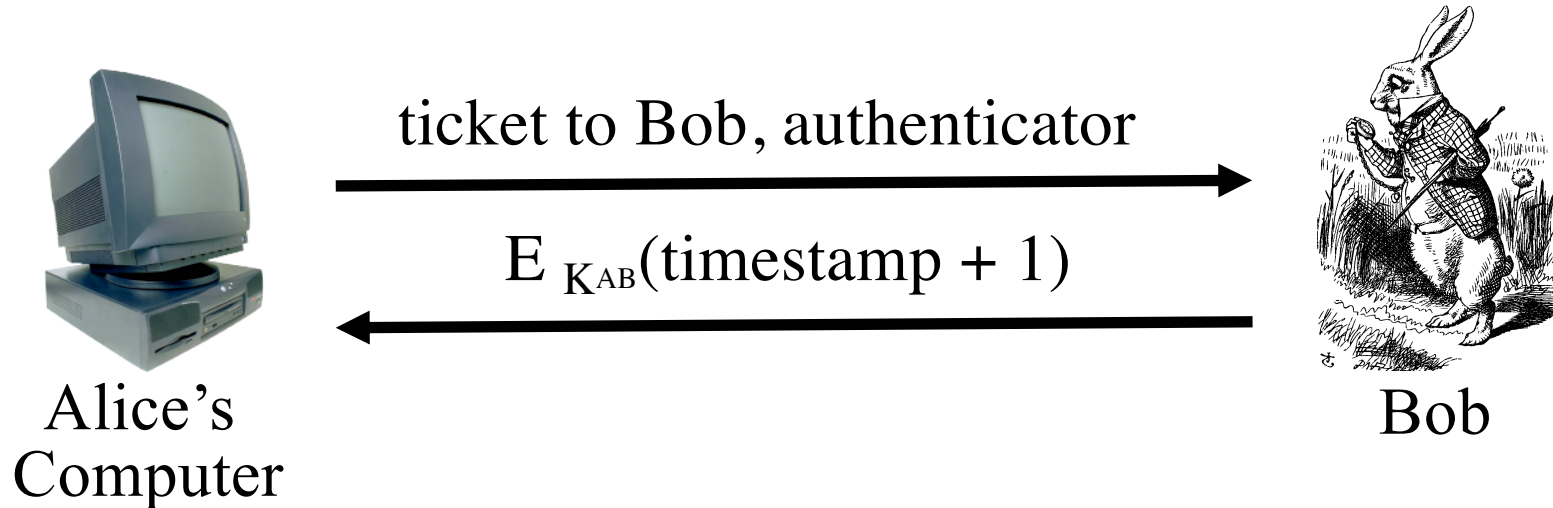
- Key $K_A = h(\text{Alice's password})$
- KDC creates session key S_A
- Alice's computer decrypts S_A and TGT
 - Then it forgets K_A
- $\text{TGT} = E_{K_{\text{KDC}}}(\text{"Alice"}, S_A)$

Alice Requests "Ticket to Bob"



- $\text{REQUEST} = (\text{TGT}, \text{authenticator})$
 - $\text{authenticator} = E_{S_A}(\text{timestamp})$
- $\text{REPLY} = E_{S_A}(\text{"Bob"}, K_{AB}, \text{ticket to Bob})$
 - $\text{ticket to Bob} = E_{K_B}(\text{"Alice"}, K_{AB})$
- KDC gets S_A from TGT to verify timestamp

Alice Uses Ticket to Bob



- ticket to Bob = $E_{K_B}(\text{"Alice"}, K_{AB})$
- authenticator = $E_{K_{AB}}(\text{timestamp})$
- Bob decrypts "ticket to Bob" to get K_{AB} which he then uses to verify timestamp

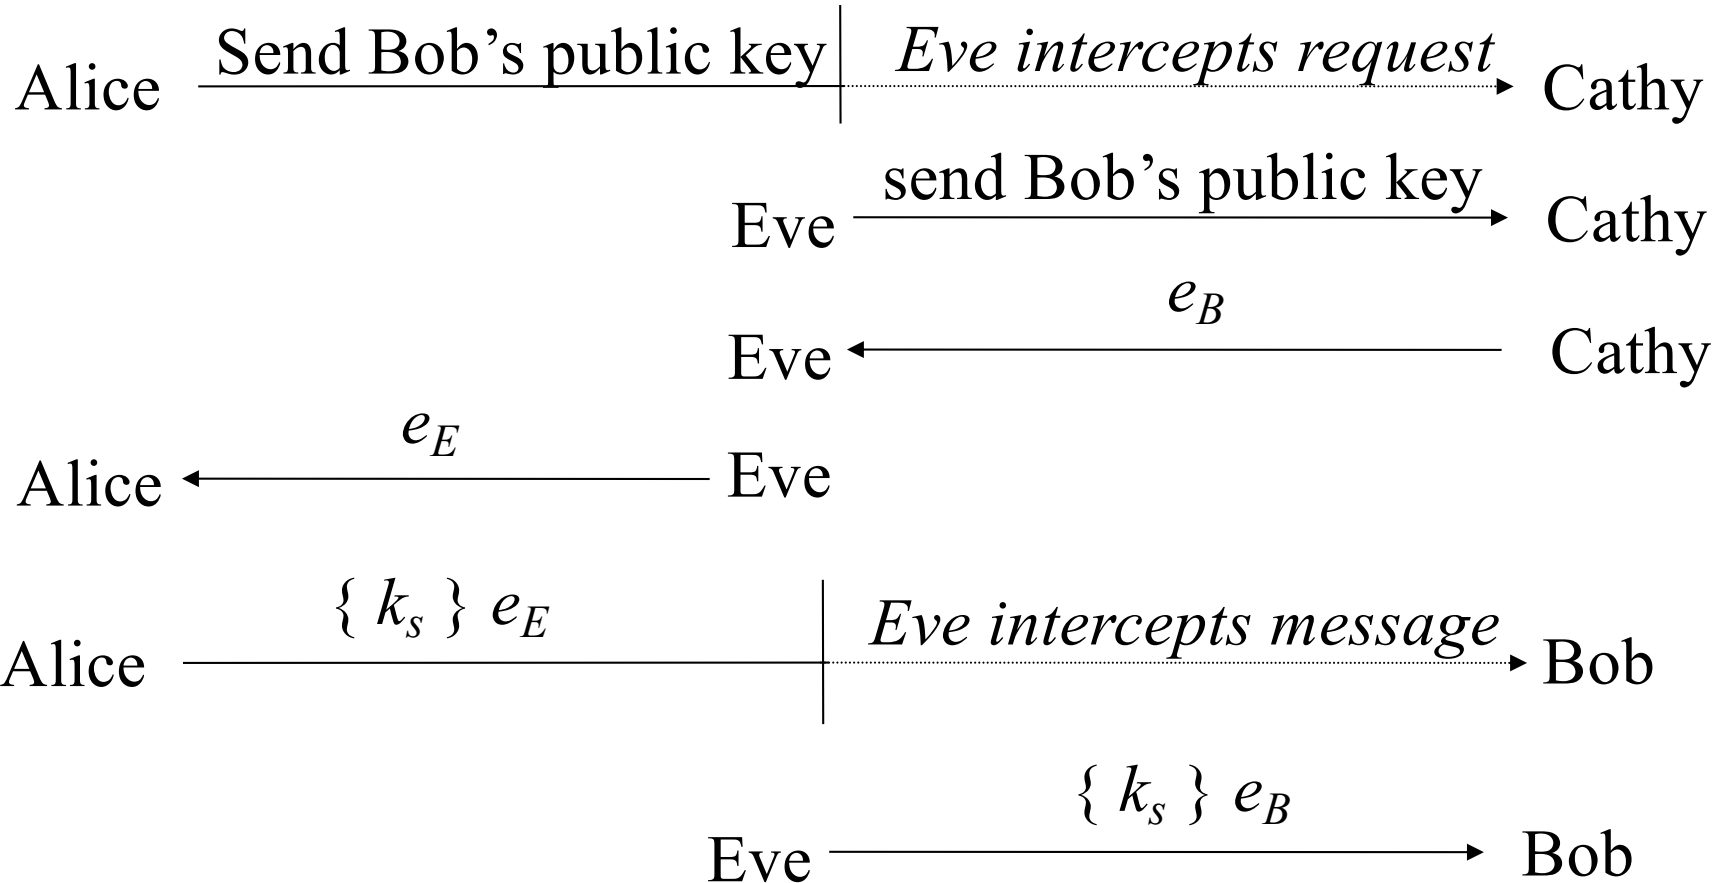

Kerberos

- Key S_A used in authentication
 - For confidentiality/integrity
- Timestamps for authentication and replay protection
- Recall, that timestamps...
 - Reduce the number of messages
 - like a nonce that is known in advance
 - But, “time” is a security-critical parameter

Kerberos Alternatives

- Could have Alice's computer remember password and use that for authentication
 - Then no KDC required
 - But hard to protect passwords
 - Also, does not scale
- Could have KDC remember session key instead of putting it in a TGT
 - Then no need for TGT
 - But **stateless** KDC is major feature of Kerberos

Man-in-the-Middle Attack



Digital Signatures



- Alice wants Bob to know that she sent message, m
 - Sends digital signature along with message
 - $m || \{h(m)\}_{d_A}$
- How would Bob verify signature?
- Could Eve intercept and change message?
- How does Bob know that Alice is the sender?

Cryptographic Key Infrastructure

- Goal: bind identity to key
- Classical: not possible as all keys are shared
 - Use protocols to agree on a shared key (see earlier)
- Public key: bind identity to public key
 - Crucial as people will use key to communicate with principal whose identity is bound to key
 - Erroneous binding means no secrecy between principals
 - Assume principal identified by an acceptable name

Certificates



- Create token (message) containing
 - Identity of principal (here, Alice)
 - Corresponding public key
 - Timestamp (when issued)
 - Other information (perhaps identity of signer)
 - Compute hash (message digest) of token

Hash encrypted by trusted authority (here, Cathy)
using private key: called a “signature”

$$C_A = e_A || \text{Alice} || T || \{h(e_A || \text{Alice} || T)\} d_C$$

X.509 Certificates

- Some certificate components in X.509v3:
 - Version
 - Serial number
 - Signature algorithm identifier: hash algorithm
 - Issuer's name; uniquely identifies issuer
 - Interval of validity
 - Subject's name; uniquely identifies subject
 - Subject's public key
 - Signature: encrypted hash

Use

- Bob gets Alice's certificate
 - If he knows Cathy's public key, he can validate the certificate
 - Decrypt the encrypted hash using Cathy's public key
 - Re-compute hash from certificate and compare
 - Check validity
 - Is the principal Alice?
 - Now Bob has Alice's public key
- Problem: Bob needs Cathy's public key to validate certificate
 - That is, **secure distribution of public keys**
 - Solution: Public Key Infrastructure (PKI) using trust anchors called Certificate Authorities (CAs) that issue certificates

PKI Trust Models

- Single Global CA
 - Unmanageable
 - Who is the universally trusted root?
- Hierarchical CA
 - Tree of CA's
 - But still need to be rooted somewhere

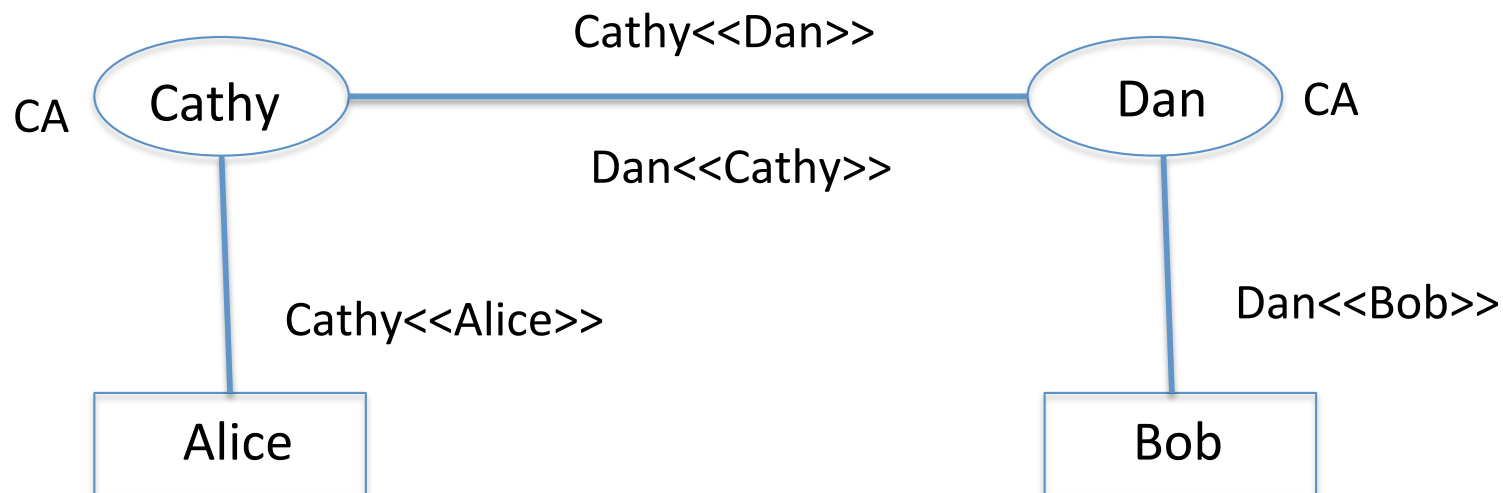
CA: Certificate Authorities

PKI Trust Models

- Hierarchical CAs with cross-certification
 - Multiple root CAs that are cross-certified
 - Cross-certification at lower levels for efficiency
- Web Model
 - Browsers come pre-configured with multiple trust anchor certificates
 - New certificates can be added
- Distributed (e.g., PGP)
 - No CA; instead, users certify each other to build a “web of trust”

Validation and Cross-Certifying

- Alice gets something signed by Bob, needs his public key
 - Certificate signed by Dan, $\text{Dan} \ll \text{Bob} \gg$
- Alice trusts Cathy who trusts Dan ($\text{Cathy} \ll \text{Dan} \gg$)
- Alice uses (known) public key of Cathy to validate $\text{Cathy} \ll \text{Dan} \gg$
- Alice uses validated public key of Dan to validate $\text{Dan} \ll \text{Bob} \gg$



Recent Root Certificate Issues

- Vast numbers of root certifiers in web browsers
- How strenuous is the background check of the certificate providers?
- How strong is the internal security of the certificate providers?
- What goes wrong with bad root certificates appear?

Key Points

- Tracking identity is important
 - Key negotiation
 - Digital signatures
- Managing the root of trust is a very hard practical problem